



# TP1 - Optimizing Memory Access

*Imad Kissami*

EL AMRANI Soufiane  
19/02/2025

## Exercise 1 :

Stride, Sum, Time (msec), Rate (MB/s)	Stride, Sum, Time (msec), Rate (MB/s)
1, 1000000.000000, 1.000000, 7629.394531	1, 1000000.000000, 1.000000, 7629.394531
2, 1000000.000000, 2.000000, 3814.697266	2, 1000000.000000, 1.000000, 7629.394531
3, 1000000.000000, 1.000000, 7629.394531	3, 1000000.000000, 2.000000, 3814.697266
4, 1000000.000000, 3.000000, 2543.131510	4, 1000000.000000, 1.000000, 7629.394531
5, 1000000.000000, 2.000000, 3814.697266	5, 1000000.000000, 2.000000, 3814.697266
6, 1000000.000000, 4.000000, 1907.348633	6, 1000000.000000, 2.000000, 3814.697266
7, 1000000.000000, 4.000000, 1907.348633	7, 1000000.000000, 3.000000, 2543.131510
8, 1000000.000000, 5.000000, 1525.878906	8, 1000000.000000, 3.000000, 2543.131510
9, 1000000.000000, 5.000000, 1525.878906	9, 1000000.000000, 3.000000, 2543.131510
10, 1000000.000000, 5.000000, 1525.878906	10, 1000000.000000, 4.000000, 1907.348633
11, 1000000.000000, 5.000000, 1525.878906	11, 1000000.000000, 4.000000, 1907.348633
12, 1000000.000000, 5.000000, 1525.878906	12, 1000000.000000, 4.000000, 1907.348633
13, 1000000.000000, 5.000000, 1525.878906	13, 1000000.000000, 5.000000, 1525.878906
14, 1000000.000000, 5.000000, 1525.878906	14, 1000000.000000, 5.000000, 1525.878906
15, 1000000.000000, 6.000000, 1271.565755	15, 1000000.000000, 5.000000, 1525.878906
16, 1000000.000000, 8.000000, 953.674316	16, 1000000.000000, 6.000000, 1271.565755
17, 1000000.000000, 5.000000, 1525.878906	17, 1000000.000000, 6.000000, 1271.565755
18, 1000000.000000, 6.000000, 1271.565755	18, 1000000.000000, 5.000000, 1525.878906
19, 1000000.000000, 6.000000, 1271.565755	19, 1000000.000000, 5.000000, 1525.878906
20, 1000000.000000, 6.000000, 1271.565755	20, 1000000.000000, 6.000000, 1271.565755

## My observations :

- Execution Time :-O2 consistently reduces execution time compared to -O0, especially for larger strides.The improvement is more pronounced for strides where cache locality is poor (e.g.,strides 8,16).
- Memory Bandwidth :-O2 improves memory bandwidth across all strides, demonstrating better utilization of CPU resources.However,the improvement diminishes for very large strides (e.g., stride16), where cache misses dominate performance.

## Exercise 2 :

```
● PS C:\Users\soufiane\OneDrive\Bureau\TP1_PL> ./mxm
Standard Matrix Multiplication:
Execution Time: 0.001000 seconds
Memory Bandwidth: 16000.00 MB/s
```

```
● PS C:\Users\soufiane\OneDrive\Bureau\TP1_PL> ./mxm

Optimized Matrix Multiplication:
Execution Time: 0.001000 seconds
Memory Bandwidth: 16000.00 MB/s
```

- Execution Time :Both versions of the matrix multiplication algorithm achieved the same execution time of 0.001000 seconds .This suggests that,for the given matrix size and hardware configuration,the reordering of loops did not result in a measurable improvement in execution time.
- Memory Bandwidth :Both versions achieved the same memory bandwidth of 16000.00 MB/s This indicates that the memory access patterns in both versions were equally efficient in terms of data transfer rates.

### Possible Reasons for Similar Results :

- ❖ The matrix size used in this experiment might have been small enough to fit entirely in the CPU cache.
- ❖ My compiler may have applied additional optimizations that mitigated the differences between the two versions.

## Exercise 3 :

```
Block Size, CPU Time (ms), Memory Bandwidth (MB/s)
8, 919.000000, 18.255948
16, 895.000000, 18.745493
32, 913.000000, 18.375921
64, 995.000000, 16.861524
128, 1154.000000, 14.538315
256, 1185.000000, 14.157988
```

- **Optimal Block Size :** From the results: The lowest CPU time is achieved with a block size of 16 , which has a CPU time of 895 ms . The highest memory bandwidth is also achieved with a block size of 16 , at 18.75 MB/s Thus, the optimal block size based on these results is 16 .
- **Reasoning for Optimal Block Size :** A block size of 16 fits well within the CPU's cache hierarchy, allowing efficient reuse of data within blocks. This minimizes cache misses and improves performance. Smaller block sizes (e.g., 8) introduce higher loop overhead due to frequent iterations over smaller chunks of data.

## Exercise 4 :

### My Analysis :

- **Memory Leak Cause :** The original program did not free the duplicate array (array\_copy), leading to a memory leak. The free\_memory function was incomplete and did not actually free the memory.
- **Fix :** I updated the free\_memory function to properly free memory. and I ensured that both array and array\_copy were freed before the program exited.